

# RELATED CONCEPTS IN COMPUTER NETWORKS

---


Some slides have been taken from: *Computer Networking: A Top Down Approach Featuring the Internet*, 3<sup>rd</sup> edition. Jim Kurose, Keith Ross. Addison-Wesley, July 2004. All material copyright 1996-2004. J.F Kurose and K.W. Ross, All Rights Reserved.

# Contents

- Computer Networks
- Internet protocol stack
- Application layer
- TCP & UDP
- Internet layer


# What's the Internet?

- *Internet*: “network of networks”
  - Interconnected ISPs
- *protocols* control sending, receiving of msgs
  - e.g., TCP, IP, HTTP, Skype, 802.11
- *Internet standards*
  - RFC: Request for comments
  - IETF: Internet Engineering Task Force



# What's the Internet?

- *Infrastructure that provides services to applications:*
  - Web, VoIP, email, games, e-commerce, social nets, ...
- *provides programming interface to apps*
  - hooks that allow sending and receiving app programs to “connect” to Internet
  - provides service options, analogous to postal service



# What' s a protocol?

## *human protocols:*

- “what' s the time?”
- “I have a question”
- introductions

... specific msgs sent


... specific actions taken  
when msgs received,  
or other events

## *network protocols:*


- machines rather than humans
- all communication activity in Internet governed by protocols

*protocols define format, order of msgs sent and received among network entities, and actions taken on msg transmission, receipt*

# Vì sao phải phân tầng?

- Đối với các hệ thống phức tạp: nguyên lý "*chia để trị*"
- Cho phép xác định rõ nhiệm vụ của mỗi bộ phận và quan hệ giữa chúng 
- Cho phép dễ dàng *bảo trì* và *nâng cấp* hệ thống
  - Thay đổi bên trong một bộ phận không ảnh hưởng đến các bộ phận khác
  - Như việc nâng cấp từ CD lên DVD player mà không phải thay loa.

# Phân tầng các chức năng hàng không



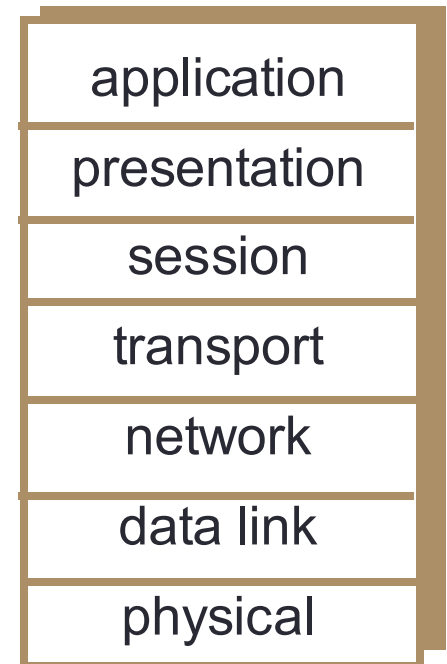
**Tầng:** Mỗi tầng có nhiệm vụ cung cấp 1 dịch vụ

- Dựa trên các chức năng của chính tầng đó
- Dựa trên các dịch vụ cung cấp bởi tầng dưới



# OSI protocol stack


- **Application layer:** defines communication between different parts of the same application
- **Presentation layer:** application data representation, data encryption, compression, conversion...
- **Session layer:** manages sessions, synchronization, recovery of data transmission process
- **Transport layer:** Transmits data between applications
- **Network layer:** Transmits data between distance network elements: Taking care of routing and forwarding data
- **Data link layer:** Transmits data between adjacent network elements.
- **Physical layer:** Transmits bits on the medium. Converting bits to physical form appropriate to the medium.






# TCP/IP protocol stack

- **application:** supporting network applications
  - FTP, SMTP, STTP
- **transport:** host-host data transfer
  - TCP, UDP
- **network:** routing of datagrams from source to destination
  - IP, routing protocols
- **link:** data transfer between neighboring network elements
  - PPP, Ethernet
- **physical:** bits “on the wire”



# OSI and TCP/IP models

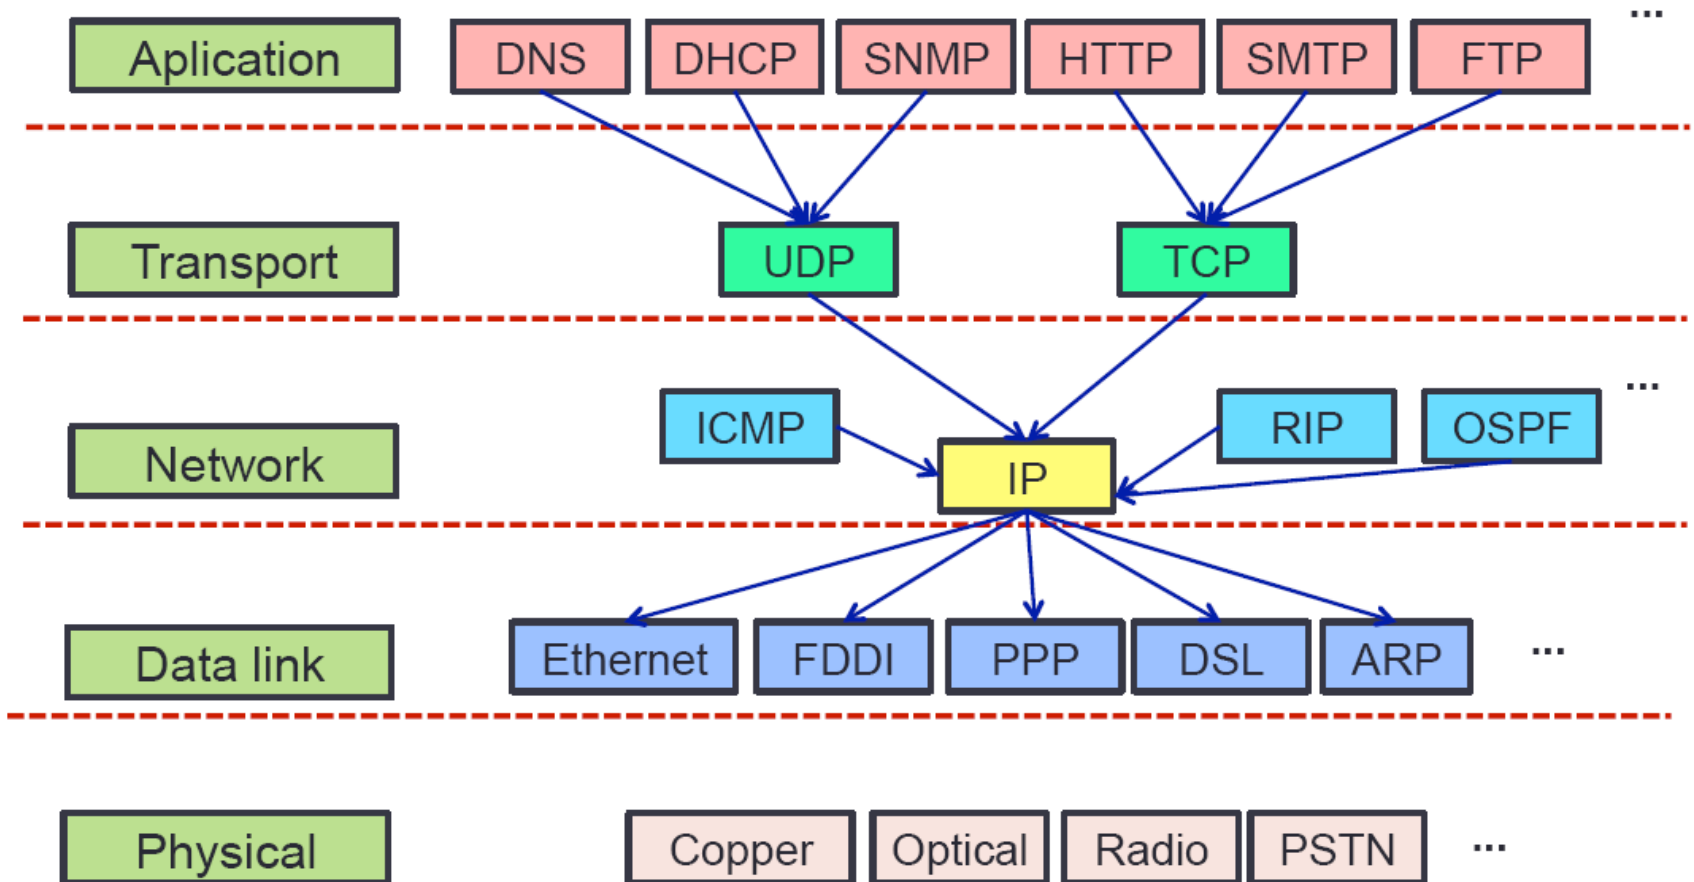


# OSI and TCP/IP models

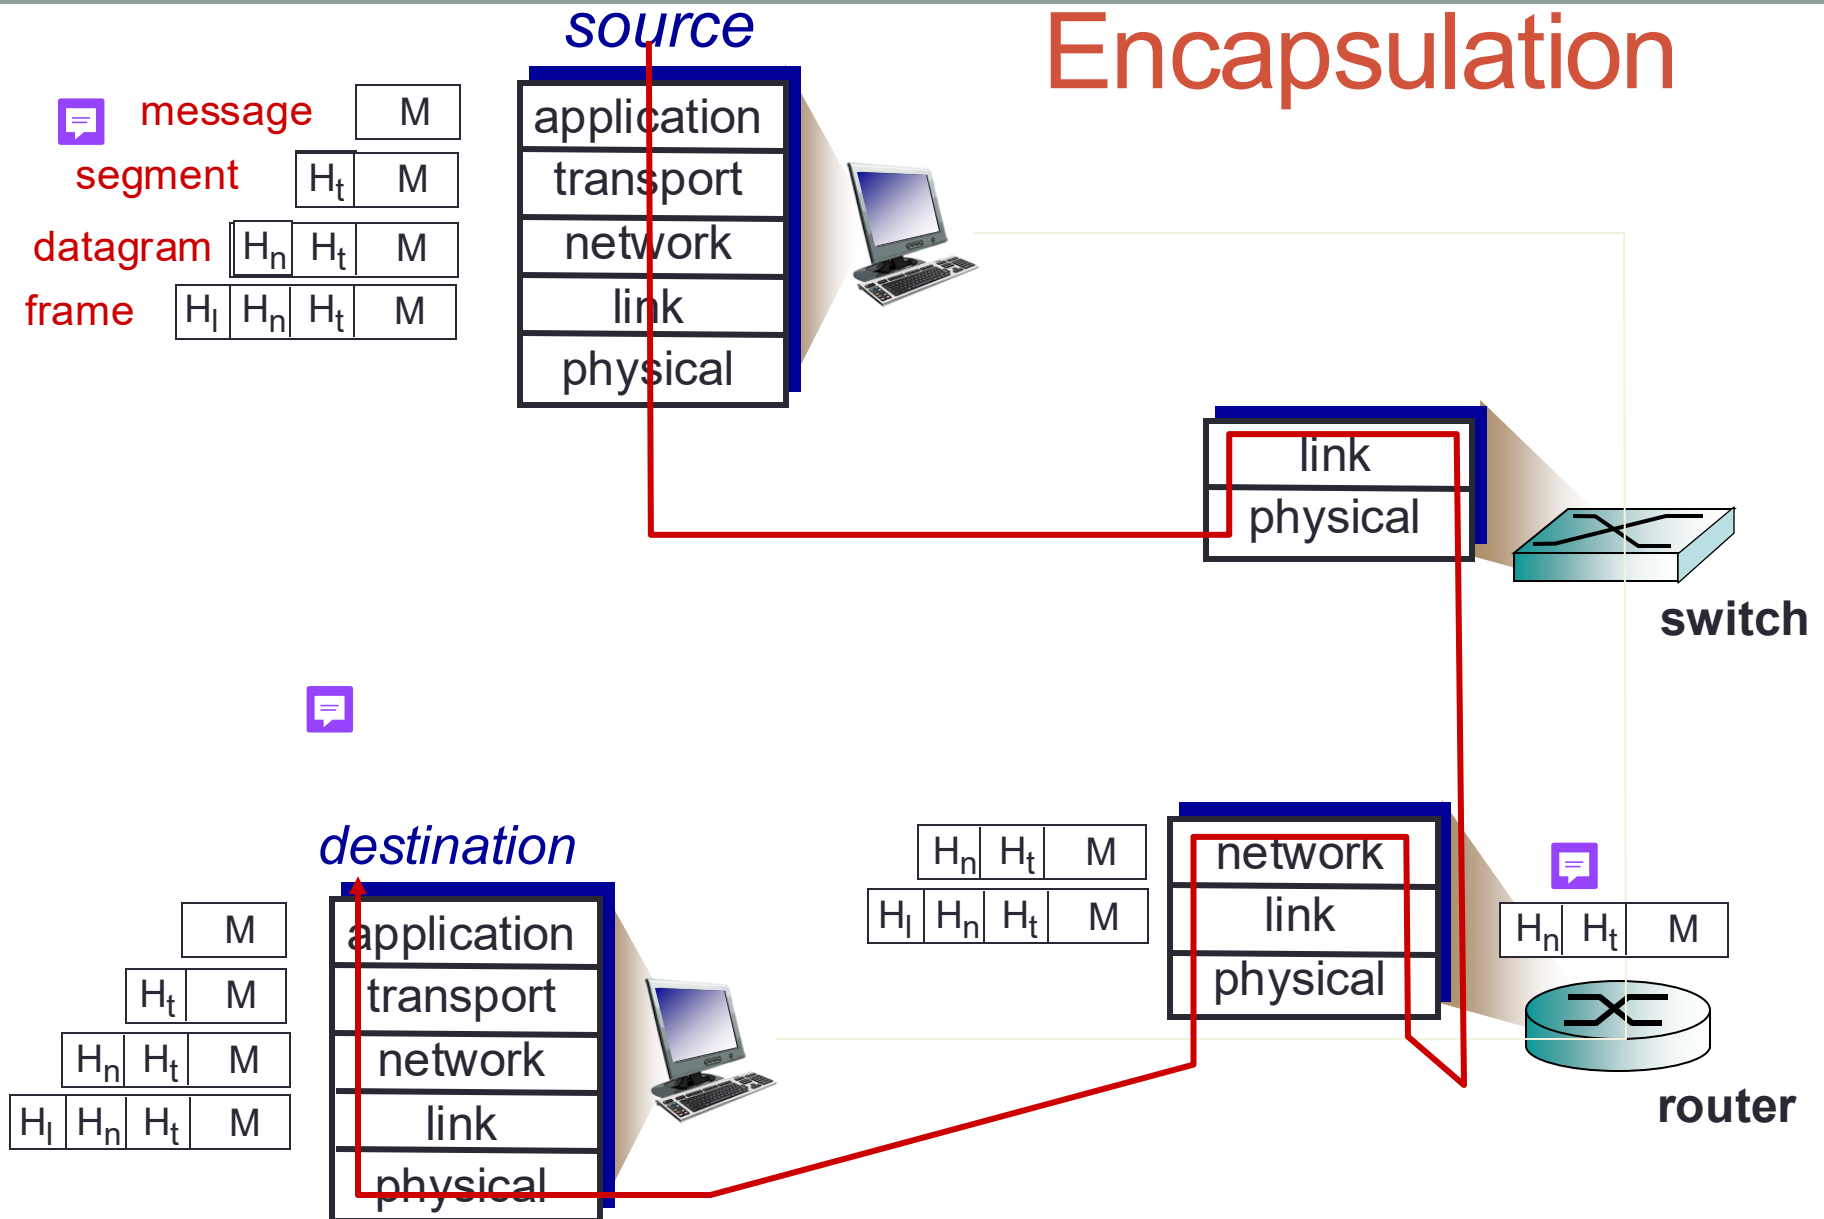
- OSI model: reference model
- TCP model: Internet model
  - Transport layer: TCP/UDP
  - Network layer: IP + routing protocols.



# Internet protocols mapping on TCP/IP



# Encapsulation



# OSI and TCP/IP models

- Layering Makes it Easier
- Application programmer
  - Doesn't need to send IP packets
  - Doesn't need to send Ethernet frames
  - Doesn't need to know how TCP implements reliability
- Only need a way to pass the data down
  - Socket is the API to access transport layer functions

# Application layer

- E-mail
- Web
- Instant messaging
- Remote login
- P2P file sharing
- Multi-user network games
- Streaming stored video clips
- Internet telephone
- Real-time video conference
- Massive parallel computing


# Creating a network app

## Write programs that

- run on different end systems and
- communicate over a network.
- e.g., Web: Web server software communicates with browser software

## No software written for devices in network core

- Network core devices do not function at app layer
- This design allows for rapid app development






# Application architectures

- Client-server
- Peer-to-peer (P2P)
- Hybrid of client-server and P2P



# Client-server architecture



## server:

- “Always” online waiting for requests from clients
- Permanent IP address

## clients:

- Request services from server
- May have dynamic IP addresses
- Do not communicate directly with each other

E.g: web, mail...


# Pure P2P architecture

- No central server
- Peers have equal role
- Peers can communicate directly to each other
- Peers do not need to be always online


Example: Gnutella

Highly scalable

But difficult to manage



# Hybrid of client-server and P2P



- Central server manages user accounts, authentication, stores data for searching process ...
- Clients communicate directly after authentication process.
- E.g. Skype
  - Server manages login process.
  - Messages, voices are transmitted directly between servers.

# Processes communicating

**Process:** program running within a host.

- within same host, two processes communicate using **inter-process communication** (defined by OS).
- processes in different hosts communicate by exchanging **messages**

**Client process:** process that initiates communication


**Server process:** process that waits to be contacted

- Note: applications with P2P architectures have client processes & server processes



# Processes communicating

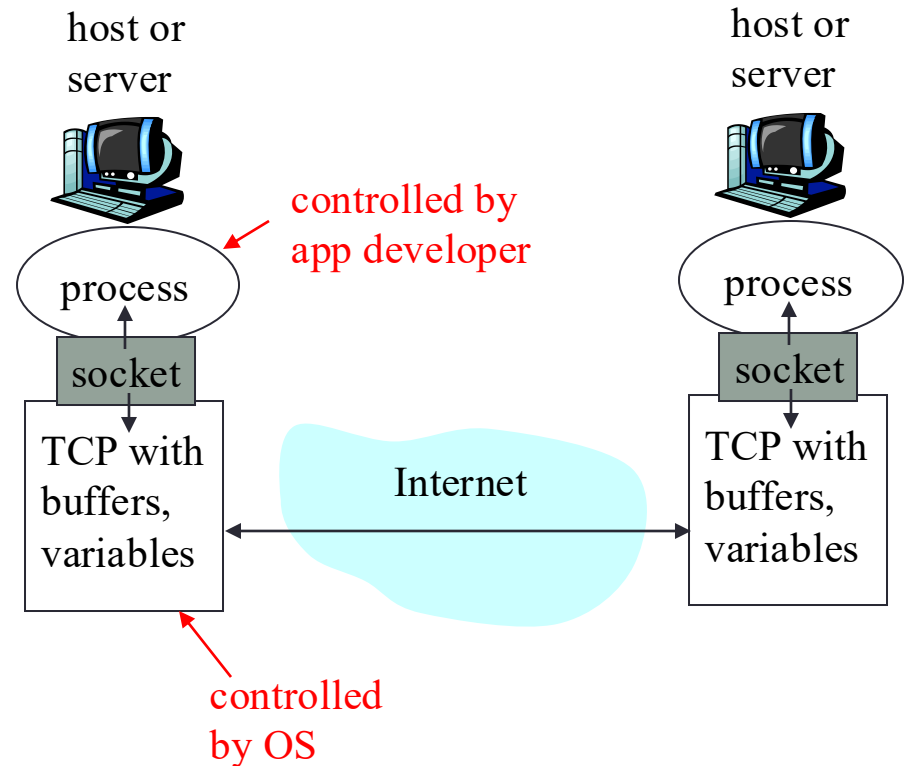
- Client process: sends request
- Server process: replies response
- Typically: single server - multiple clients
- The server does not need to know anything about the client
- The client should always know something about the server
  - at least the socket address of the server



# Sockets

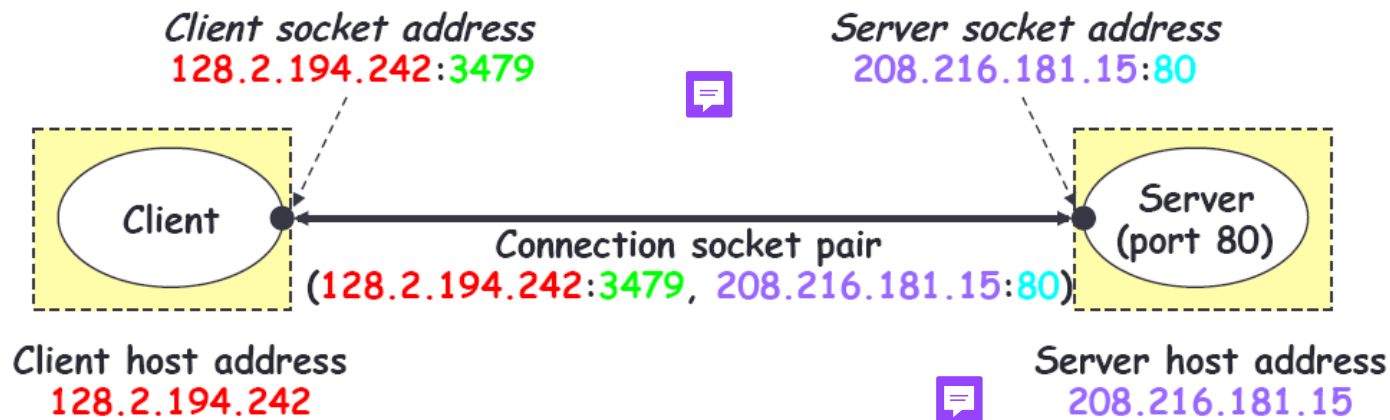
- process sends/receives messages to/from its **socket**
- Defined by
  - Port number
  - IP Address
  - TCP/UDP

} Socket  
address



# Internet connections (TCP/IP)

- Address the machine on the network
  - By IP address
- Address the process/application
  - By the “port”-number
- The pair of *IP-address* + *port* – makes up a “socket-address”



*Note: 3479 is an ephemeral port allocated by the kernel*

*Note: 80 is a well-known port associated with Web servers*



# Internet connections (TCP/IP)

- Need to open two sockets of both sides
  - Client socket
  - Server socket
- Client application send/receive data to server through client socket
- Server application send/receive data to client through client socket
- Make two sockets talk to each other.

# App-layer protocol defines

- Types of messages exchanged, eg, request & response messages
- Syntax of message types: what fields in messages & how fields are delineated
- Semantics of the fields, ie, meaning of information in fields
- Rules for when and how processes send & respond to messages

# What transport service does an app need?

## Data loss

- some apps (e.g., audio) can tolerate some loss
- other apps (e.g., file transfer, telnet) require 100% reliable data transfer


## Timing


- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

## Bandwidth

- ❑ some apps (e.g., multimedia) require minimum amount of bandwidth to be “effective”
- ❑ other apps (“elastic apps”) make use of whatever bandwidth they get

# Transport services and protocols

- provide *logical communication* between app processes running on different hosts
- transport protocols run in end systems
  - send side: breaks app messages into **segments**, passes to network layer
  - rcv side: reassembles segments into messages,  passes to app layer
- more than one transport protocol available to apps
  - Internet: TCP and UDP



# Internet transport protocols services

## *TCP service:*

- *reliable transport* between sending and receiving process
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *does not provide*: timing, minimum throughput guarantee, security
- *connection-oriented*: setup required between client and server processes

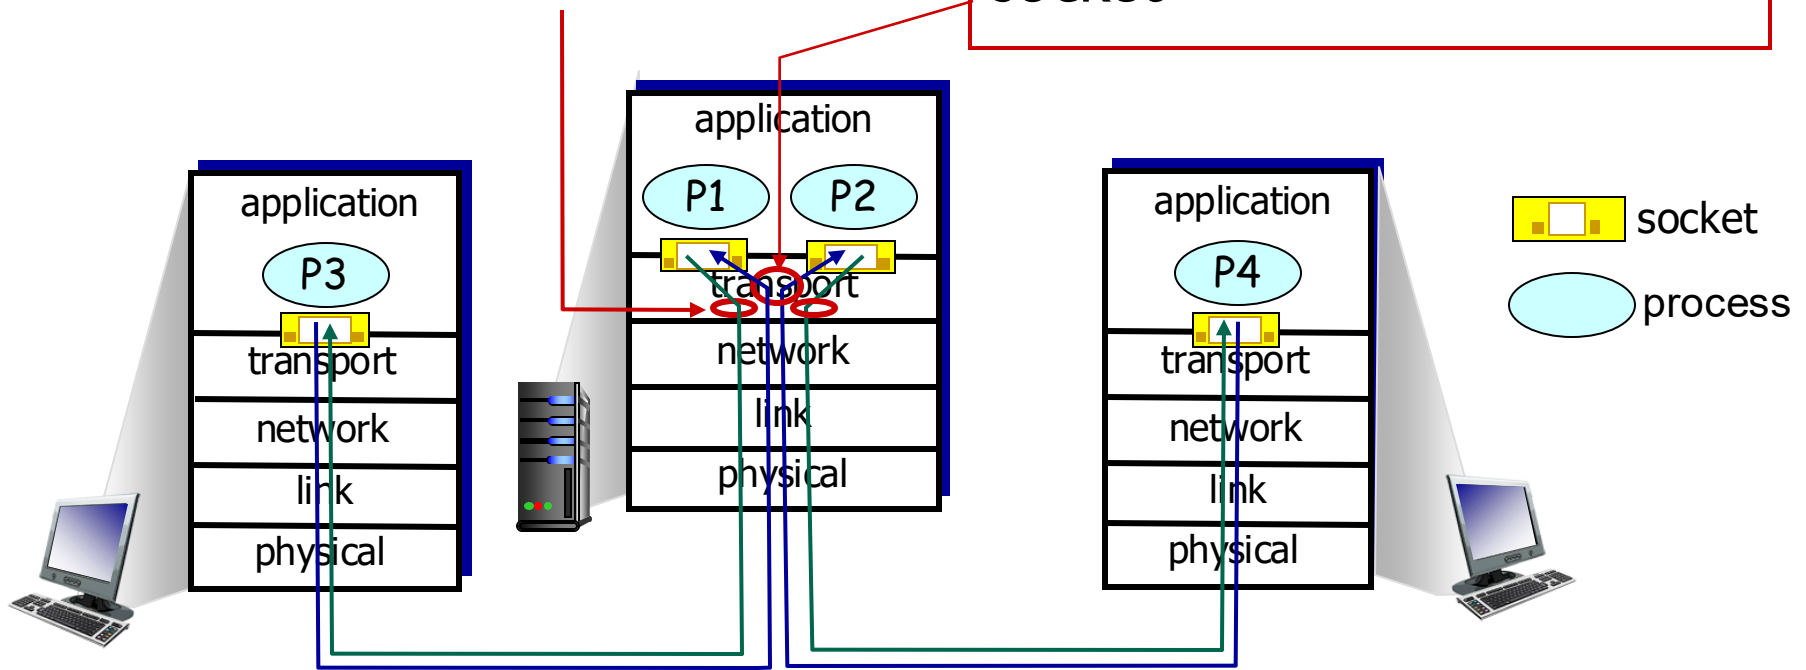
## *UDP service:*

- *unreliable data transfer* between sending and receiving process
- *does not provide*: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

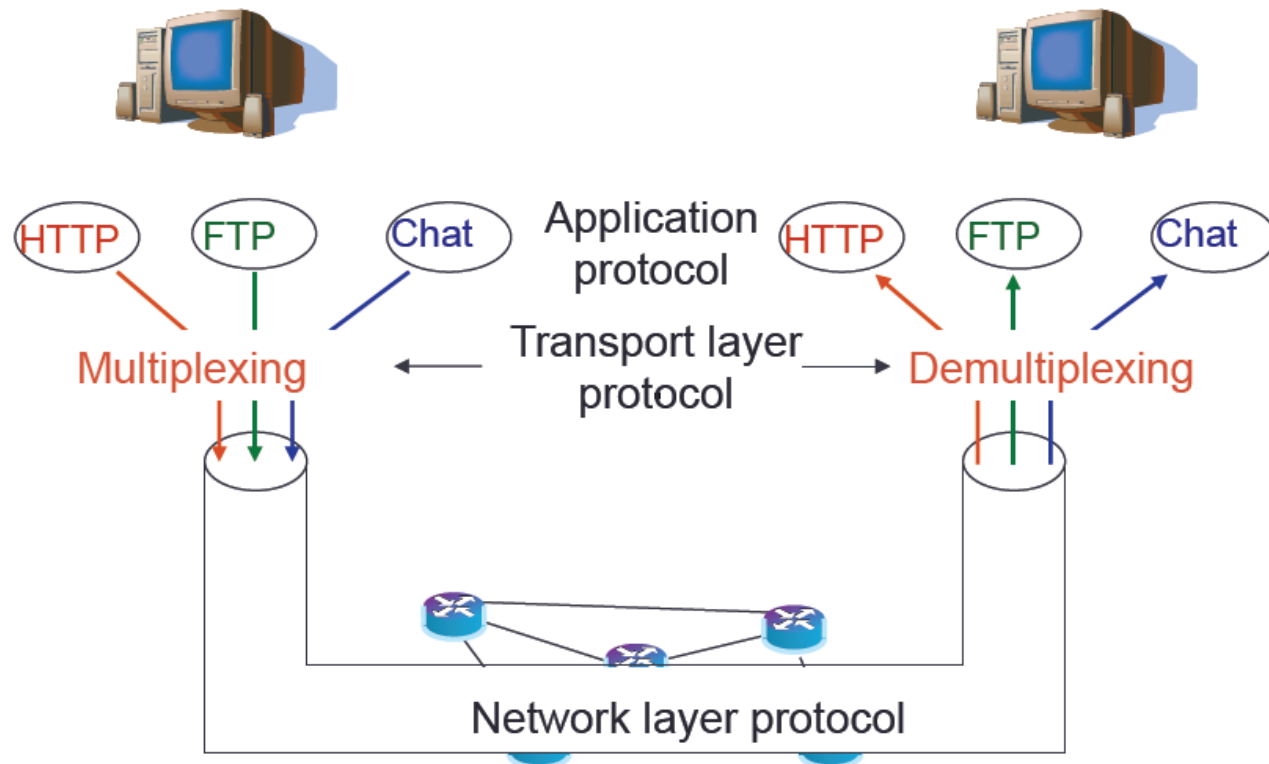
# Multiplexing/demultiplexing

*multiplexing at sender:*  
handle data from multiple sockets, add transport header (later used for demultiplexing)

*demultiplexing at receiver:*  
use header info to deliver received segments to correct socket



# Transport layer Mux/Demux

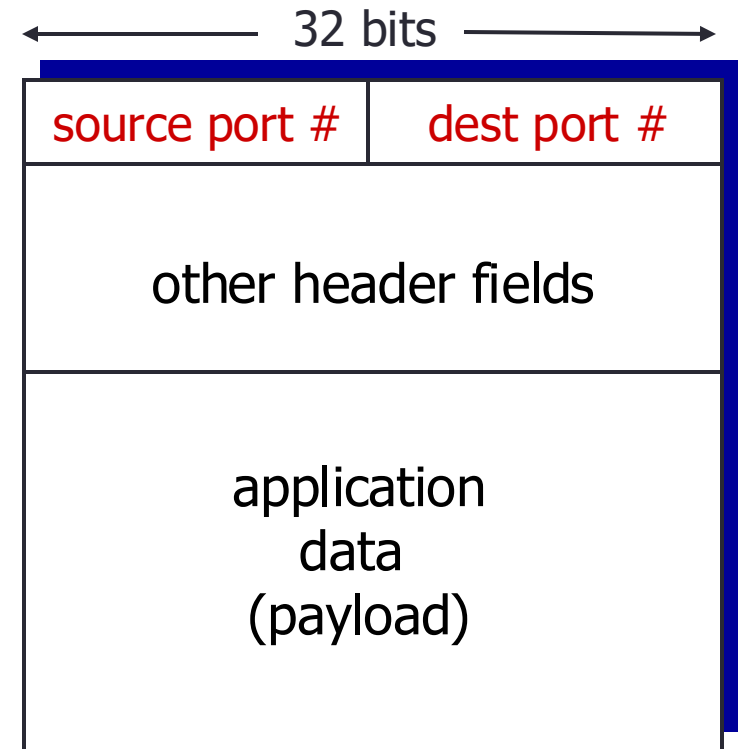


# How demultiplexing works

## ❖ host receives IP datagrams

- each datagram has source IP address, destination IP address
- each datagram carries one transport-layer segment
- each segment has source, destination port number

## ❖ host uses *IP addresses* & *port numbers* to direct segment to appropriate socket



TCP/UDP segment format



# UDP: User Datagram Protocol [RFC 768]

- “best effort” service, UDP segments may be:
  - lost
  - delivered out of order to app
- *connectionless*:
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others

## Why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small segment header
- no congestion control: UDP can blast away as fast as desired

# UDP demultiplexing

- Create sockets with port numbers:

```
mySocket = socket(AF_INET,  
    SOCK_DGRAM, 0)
```

- UDP socket identified by two-tuple:

(dest IP address, dest port number)

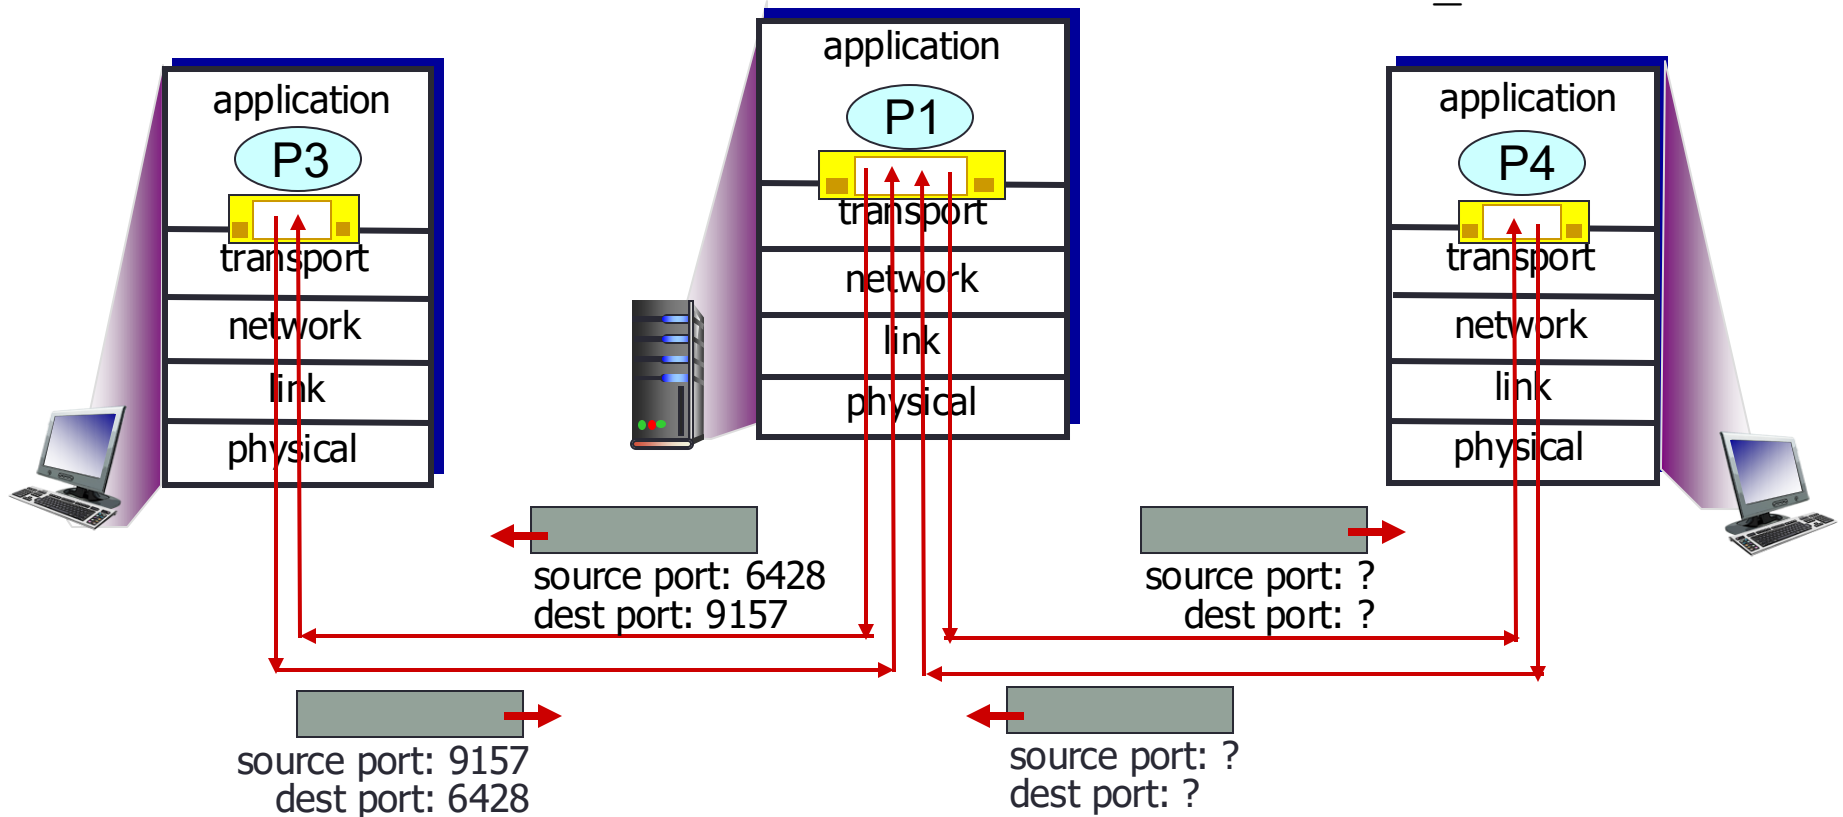
- When host receives UDP segment:
  - checks destination port number in segment
  - directs UDP segment to socket with that port number
- IP datagrams with different source IP addresses and/or source port numbers directed to same socket

# UDP demux

```
mySocket =  
  socket(AF_INET,  
    SOCK_DGRAM, 0);
```

```
serverSocket =  
  socket(AF_INET,  
    SOCK_DGRAM, 0);  
  bind(...)
```


```
mySocket =  
  socket(AF_INET,  
    SOCK_DGRAM, 0);
```



# TCP: Overview


RFCs: 793, 1122, 1323, 2018, 2581

- **point-to-point:**
  - one sender, one receiver
- **reliable, in-order *byte stream*:**
  - no “message boundaries”
- **pipelined:**
  - TCP congestion and flow control set window size
- ***send & receive buffers***
- **full duplex data:**
  - bi-directional data flow in same connection
  - MSS: maximum segment size
- **connection-oriented:**
  - handshaking (exchange of control msgs) init's sender, receiver state before data exchange
- **flow controlled:**
  - sender will not overwhelm receiver



# TCP Connection Management: Setup

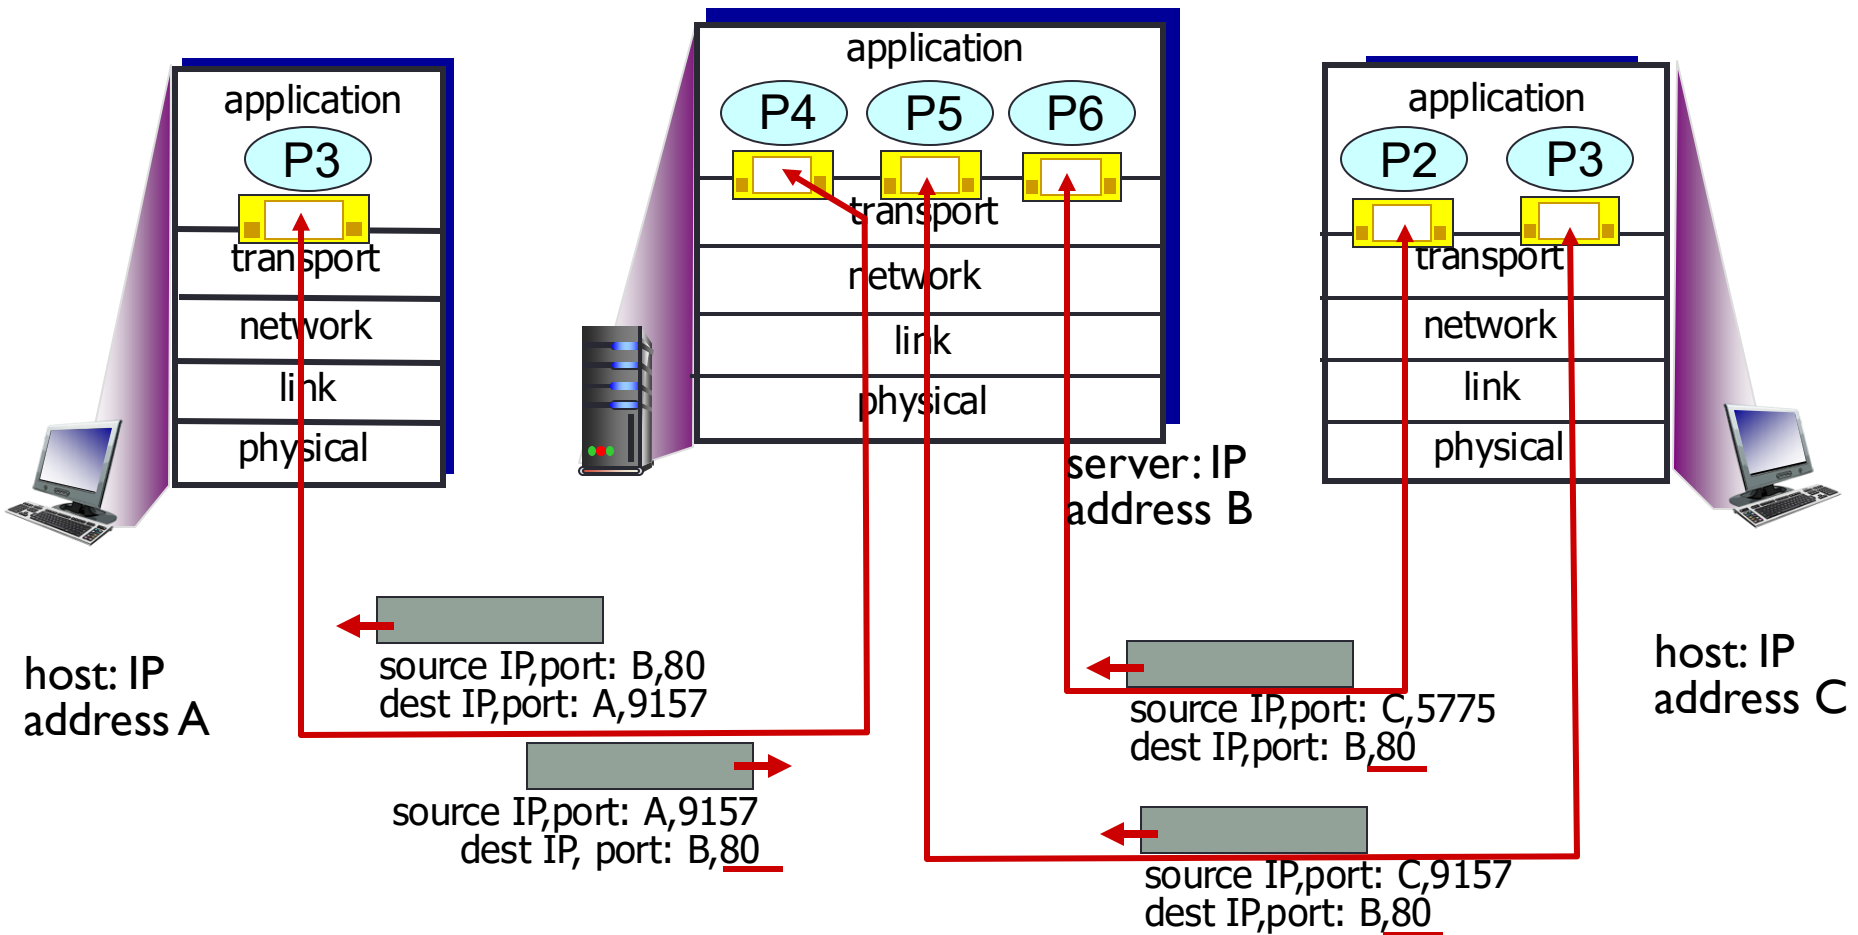
- Connection oriented protocol
  - 3-step connection opening
- Reliable protocol
  - Re-transmission on error
- Flow control
- Congestion control



# Connection-oriented demux

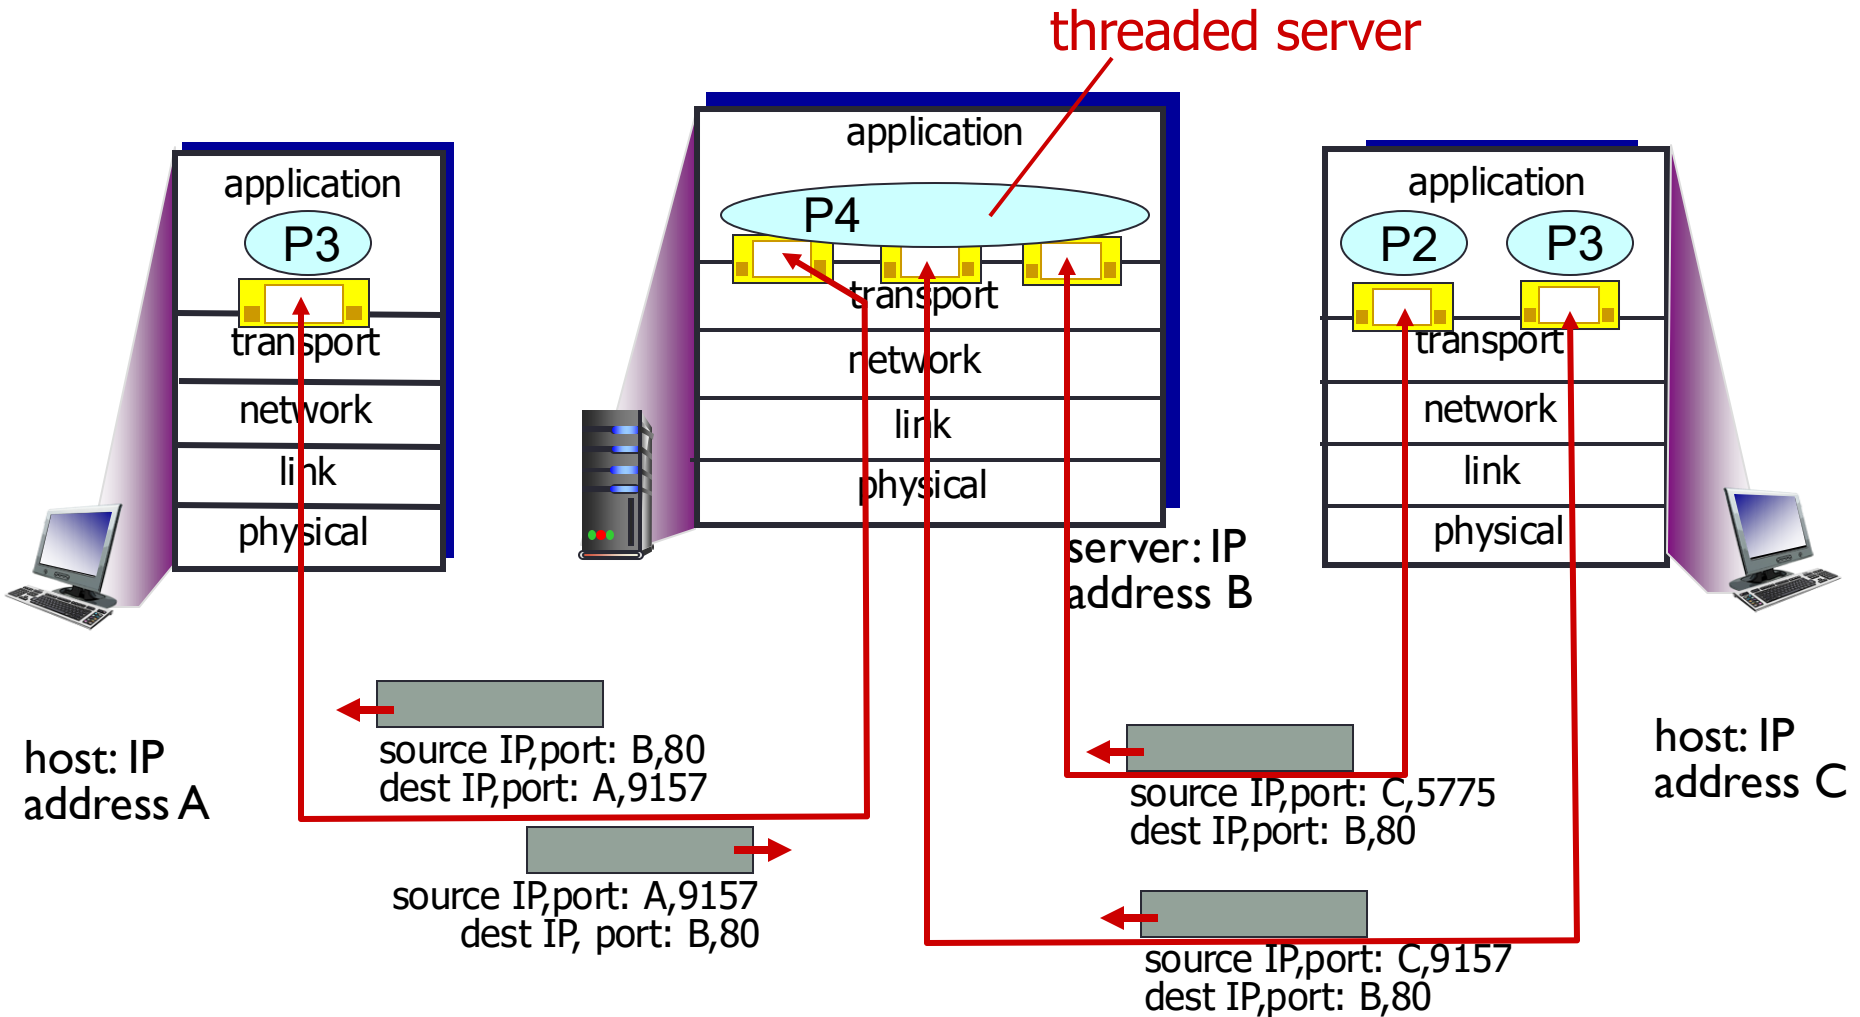
- TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- recv host uses all four values to direct segment to appropriate socket
- Server host may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
- Web servers have different sockets for each connecting client

# Connection-oriented demux: example



three segments, all destined to IP address: B,  
dest port: 80 are demultiplexed to *different* sockets

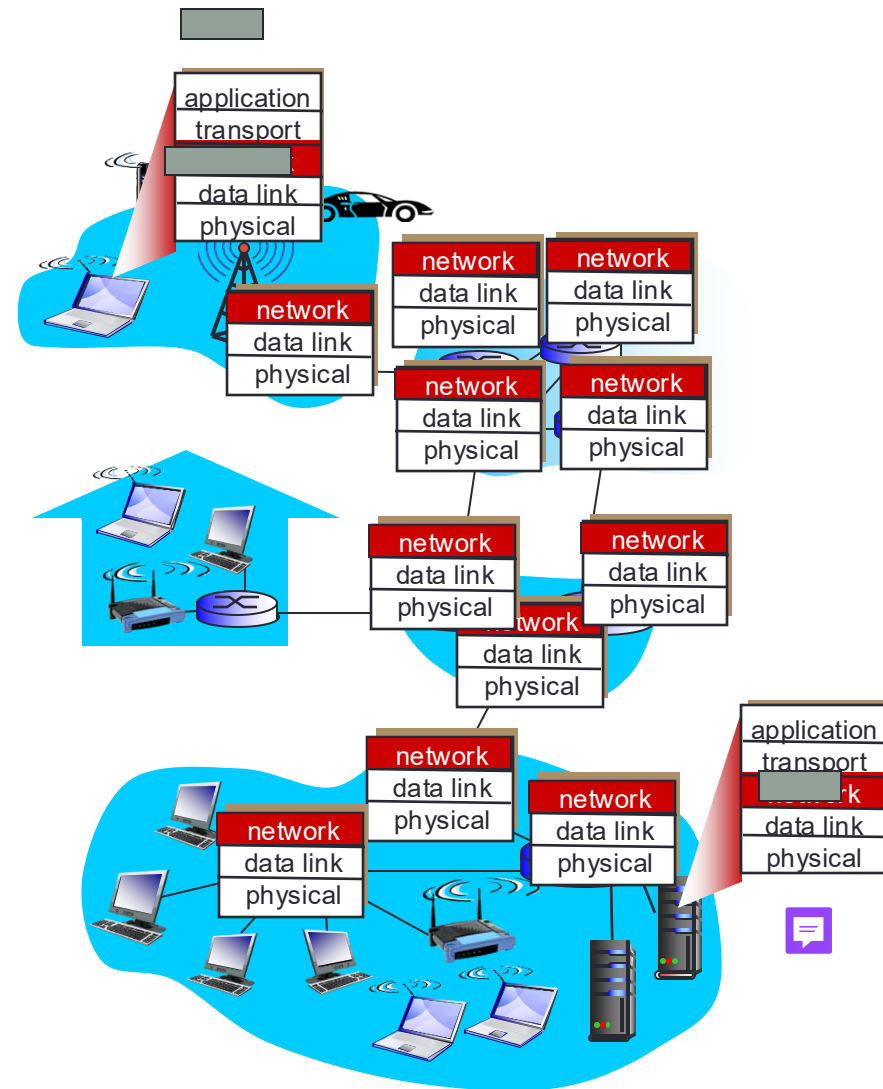
# Connection-oriented demux: example





# Network layer

- transport segment from sending to receiving host
- on sending side encapsulates segments into datagrams
- on receiving side, delivers segments to transport layer
- network layer protocols in *every* host, router
- router examines header fields in all IP datagrams passing through it



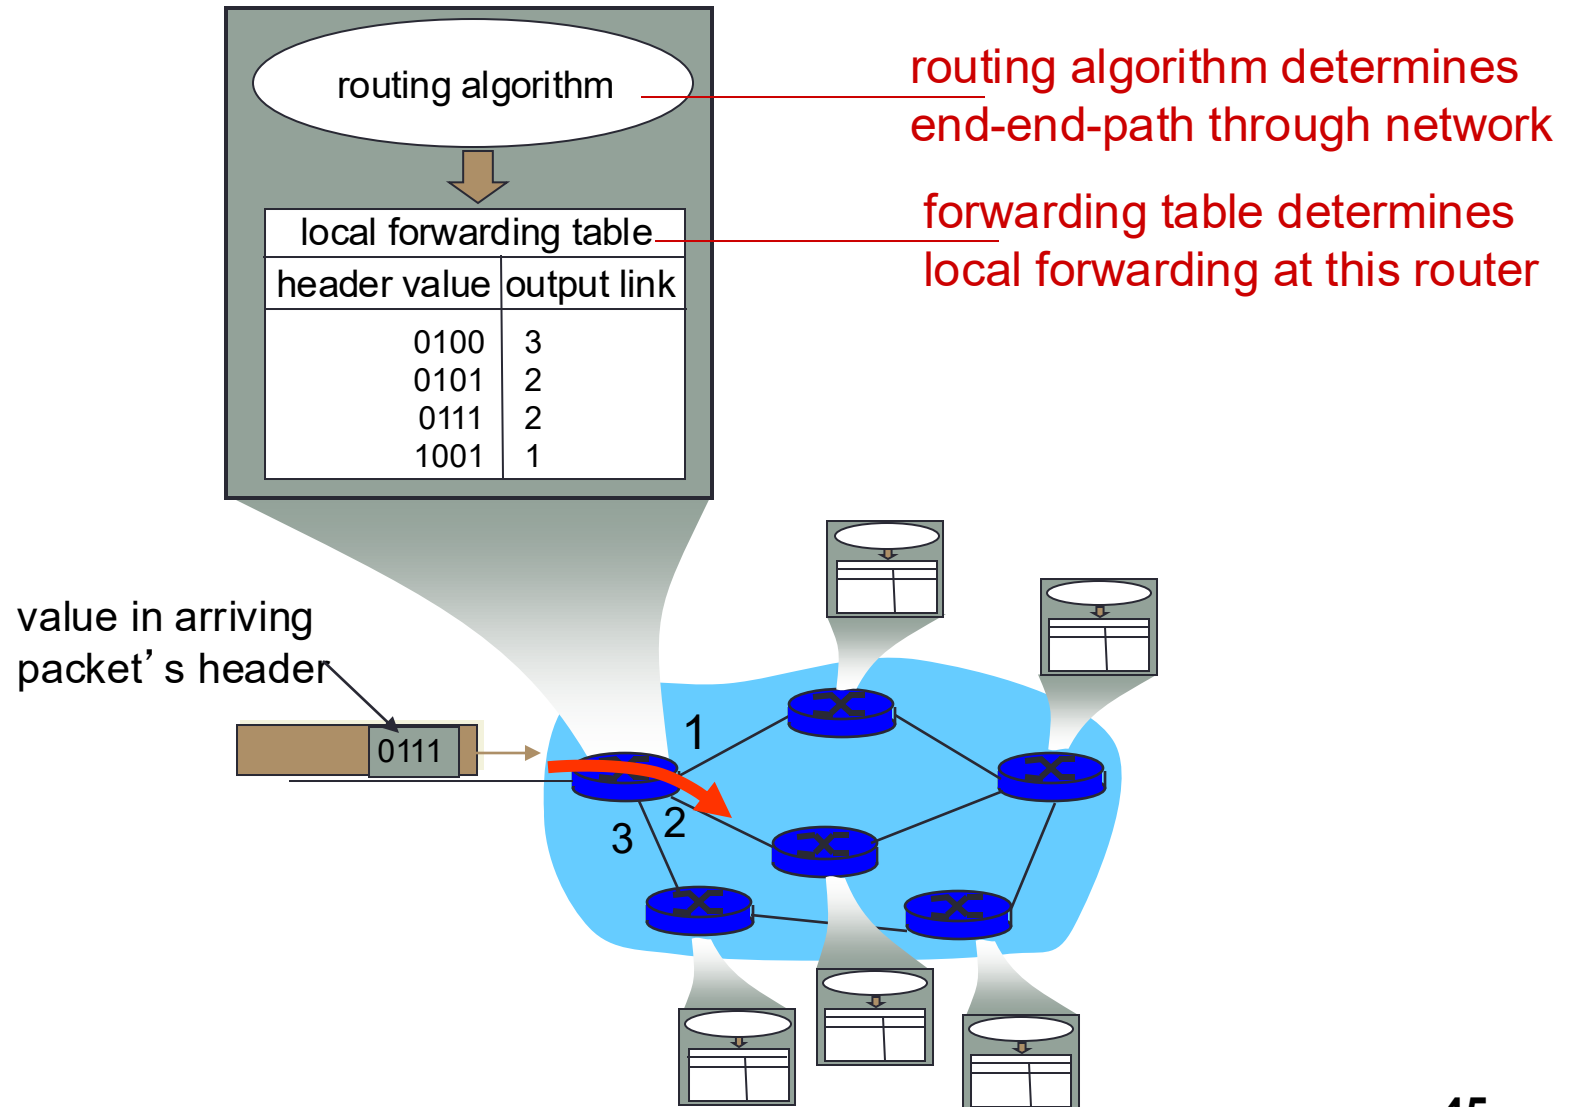
# Two key network-layer functions

- *forwarding*: move packets from router's input to appropriate router output
- *routing*: determine route taken by packets from source to dest.
  - *routing algorithms*

## *analogy:*


- *routing*: process of planning trip from source to dest
- *forwarding*: process of getting through single interchange

# Interplay between routing and forwarding



# Why an internet layer?


- ❑ Why not one big flat LAN?
  - Different LAN protocols
  - Flat address space not scalable
- ❑ IP provides:
  - Global addressing
  - Scaling to WANs
  - Virtualization of network isolates end-to-end protocols from network details/changes



"hourglass model"  
(Steve Deering)

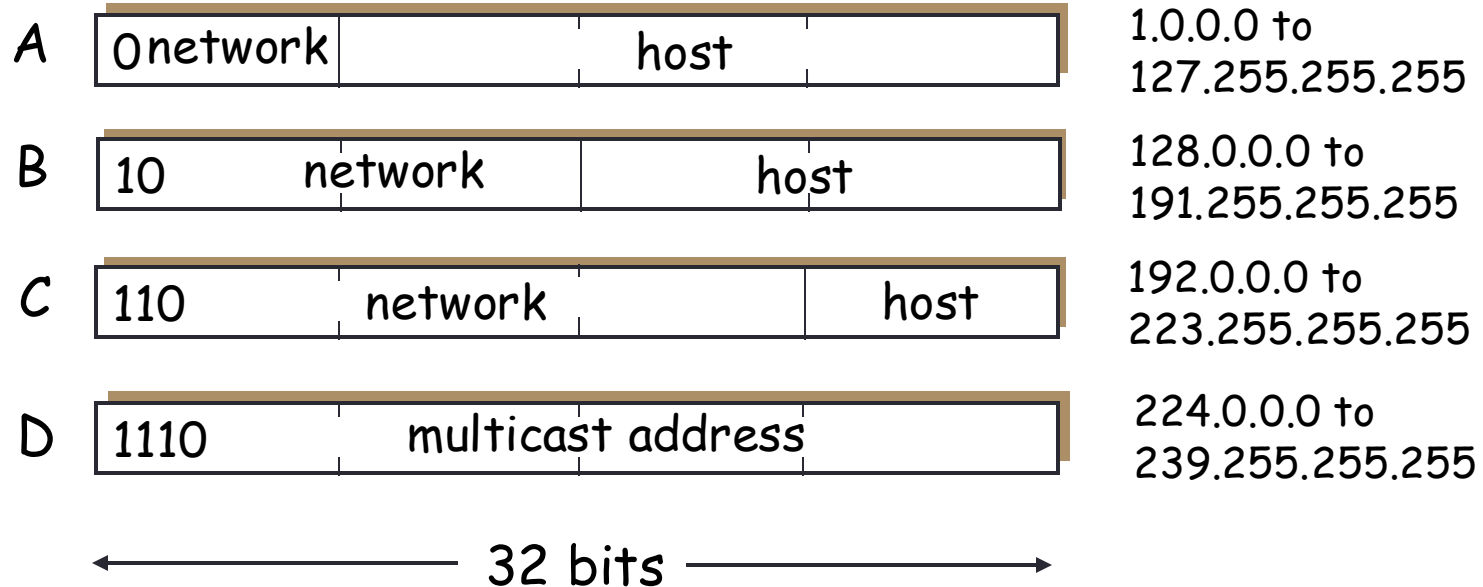
# IP addressing: introduction

- **IP address:** 32-bit identifier for host, router interface
- **interface:** connection between host/router and physical link
  - router's typically have multiple interfaces
  - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)
- **IP addresses associated with each interface**



# IP addressing: “class-full”

class




- Classful addressing:
  - inefficient use of address space, address space exhaustion
  - e.g., class B net allocated enough addresses for 65K hosts, even if only 2K hosts in that network

# IP addressing: “class-less”

## CIDR: Classless InterDomain Routing

- subnet portion of address of arbitrary length
- address format: **a.b.c.d/x**, where x is # bits in subnet portion of address



# Address Allocation for Private Internets

- RFC1918

Private address	10.0.0.0/8
	172.16.0.0/16 → 172.31.0.0/16
	192.168.0.0/24 → 192.168.255.0 /24
Loopback address	127.0.0.0 /8
Multicast address	224.0.0.0
	~239.255.255.255

- Link local address: 169.254.0.0/16




# IP

## IP: Internet Protocol

- Forward data packet between distance network nodes (routers or hosts)
- Using routing table built by routing protocols such as OSPF, RIP ...
- IP address
  - Is assigned to each network interface
  - IP v4: 32 bits
    - 133.113.215.10
  - IP v6: 128 bits
    - 2001:200:0:8803::53
- A host may have a domain name
  - Conversion IP <--> domain name: DNS
  - Ex: soict.hust.edu.vn <--> 202.191.56.68

# NAT: network address translation



*all* datagrams *leaving* local network have *same* single source NAT IP address: 138.76.29.7, different source port numbers

datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)

# NAT: network address translation

*motivation:* local network uses just one IP address as far as outside world is concerned:

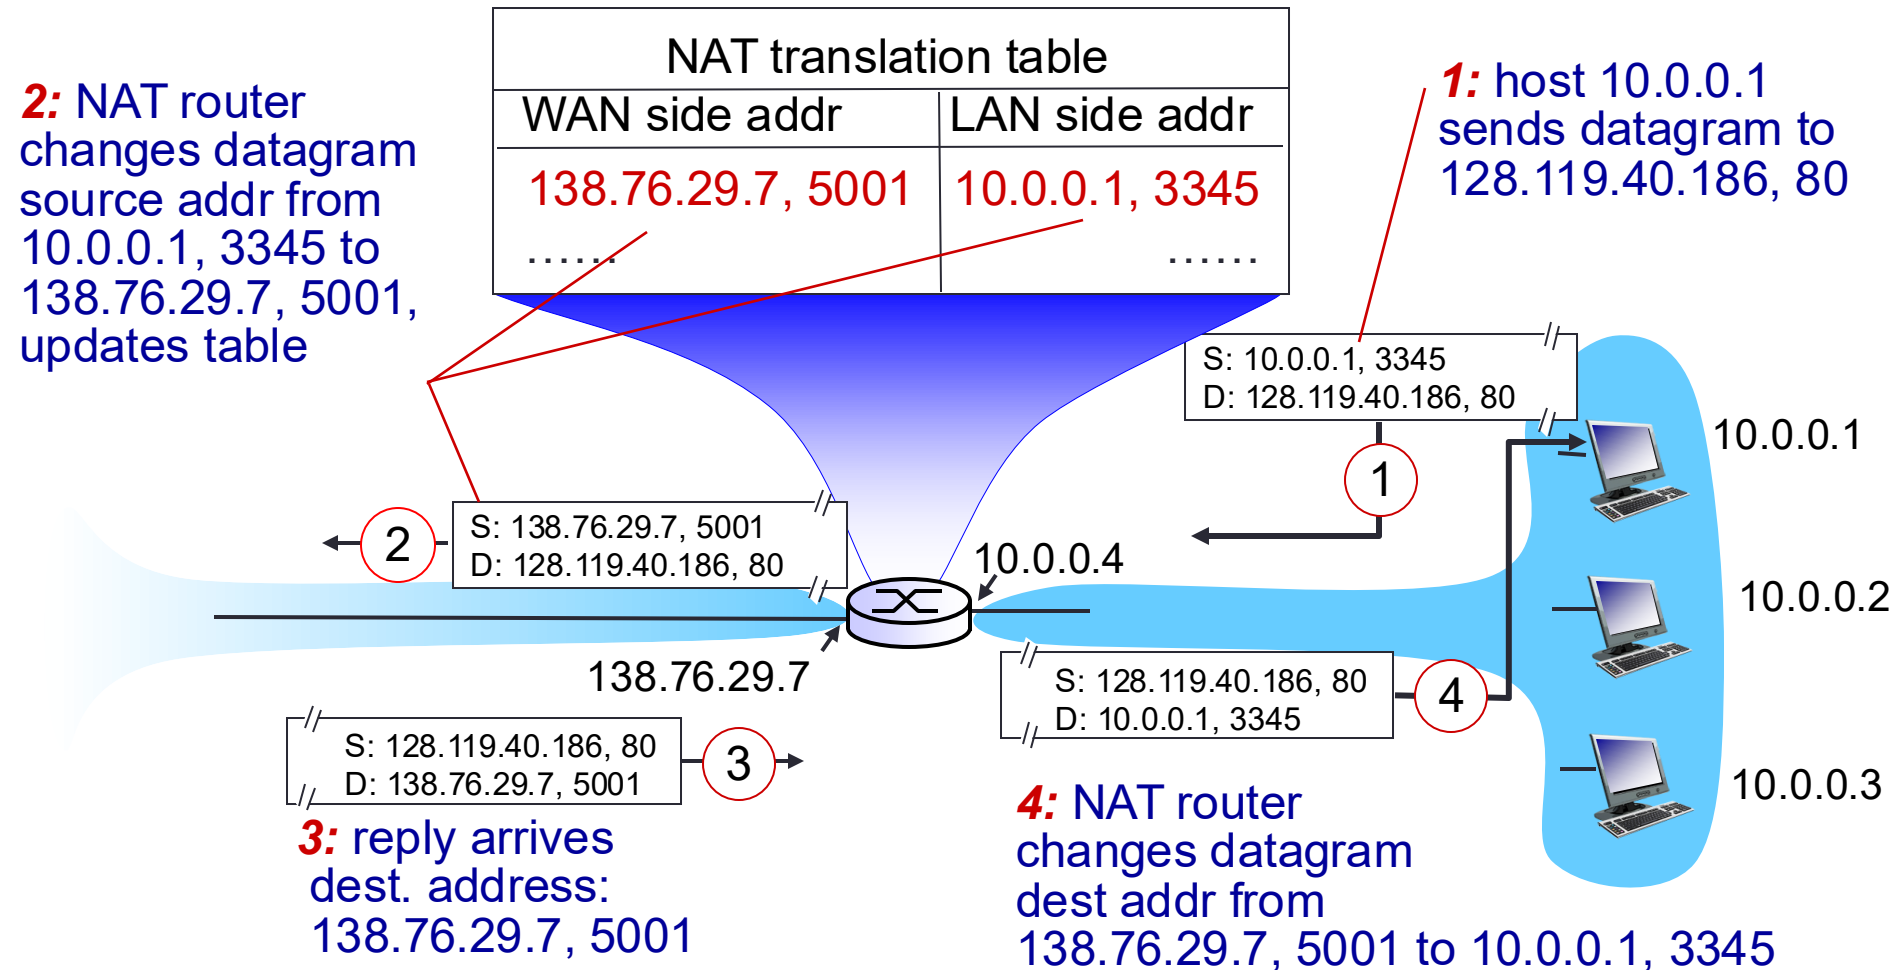
- range of addresses not needed from ISP: just one IP address for all devices
- can change addresses of devices in local network without notifying outside world
- can change ISP without changing addresses of devices in local network
- devices inside local net not explicitly addressable, visible by outside world (a security plus)

# NAT: network address translation


*implementation:* NAT router must:

- *outgoing datagrams: replace* (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)  
... remote clients/servers will respond using (NAT IP address, new port #) as destination addr
- *remember (in NAT translation table)* every (source IP address, port #) to (NAT IP address, new port #) translation pair
- *incoming datagrams: replace* (NAT IP address, new port #) in dest fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

# NAT: network address translation




# NAT: network address translation

- 16-bit port-number field: 
  - 60,000 simultaneous connections with a single LAN-side address!
- NAT is controversial:
  - routers should only process up to layer 3
  - violates end-to-end argument
    - NAT possibility must be taken into account by app designers, e.g., P2P applications
  - address shortage should instead be solved by IPv6


# NAT traversal problem

- client wants to connect to server with address 10.0.0.1
  - server address 10.0.0.1 local to LAN (client can't use it as destination addr)
  - only one externally visible NATed address: 138.76.29.7
- **solution1:** statically configure NAT to forward incoming connection requests at given port to server
  - e.g., (123.76.29.7, port 2500) always forwarded to 10.0.0.1 port 25000



# NAT traversal problem

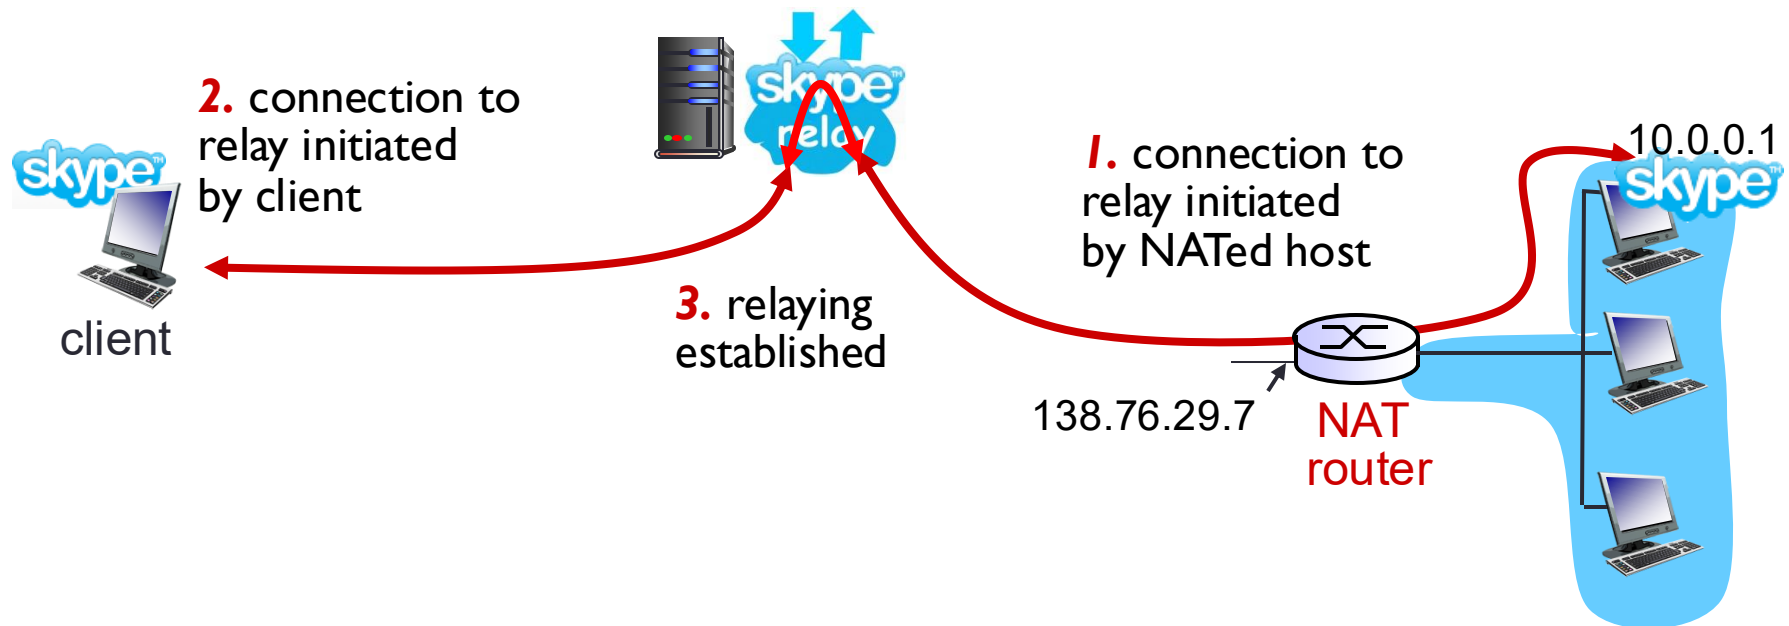
- *solution 2*: Universal Plug and Play (UPnP) Internet Gateway Device (IGD) Protocol. Allows NATed host to:
    - ❖ learn public IP address (138.76.29.7)
    - ❖ add/remove port mappings (with lease times)
- i.e., automate static NAT port map configuration





# NAT traversal problem

- **solution 3:** relaying (used in Skype)
  - NATed client establishes connection to relay
  - external client connects to relay
  - relay bridges packets between to connections



# NAT traversal problem

- *solution 4*: NAT hole punching. Example: STUN protocol

