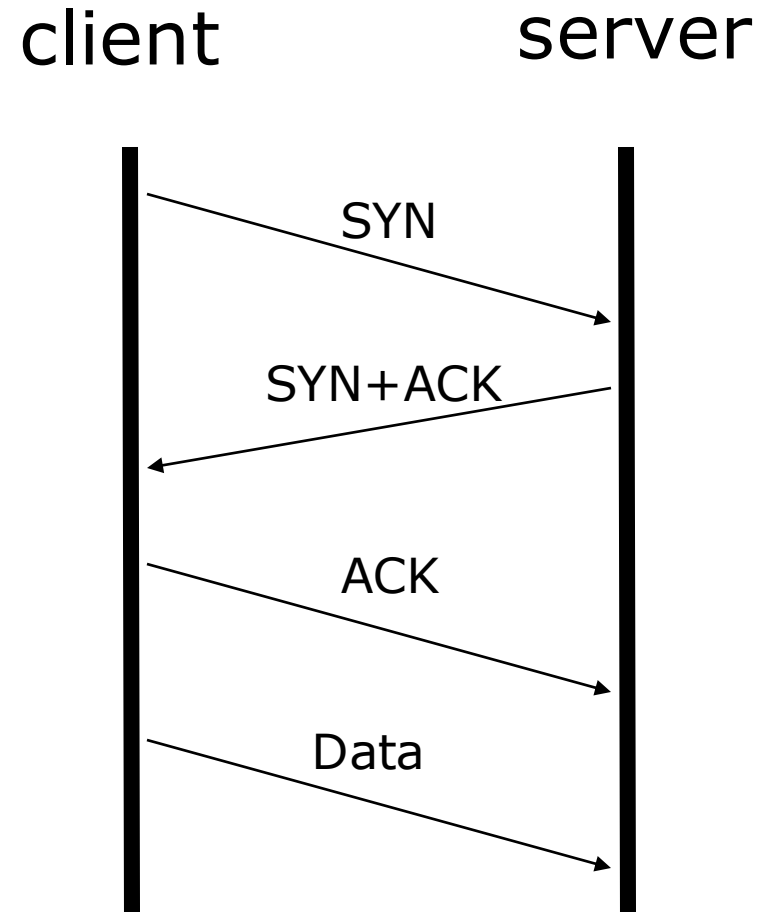


TCP SOCKETS

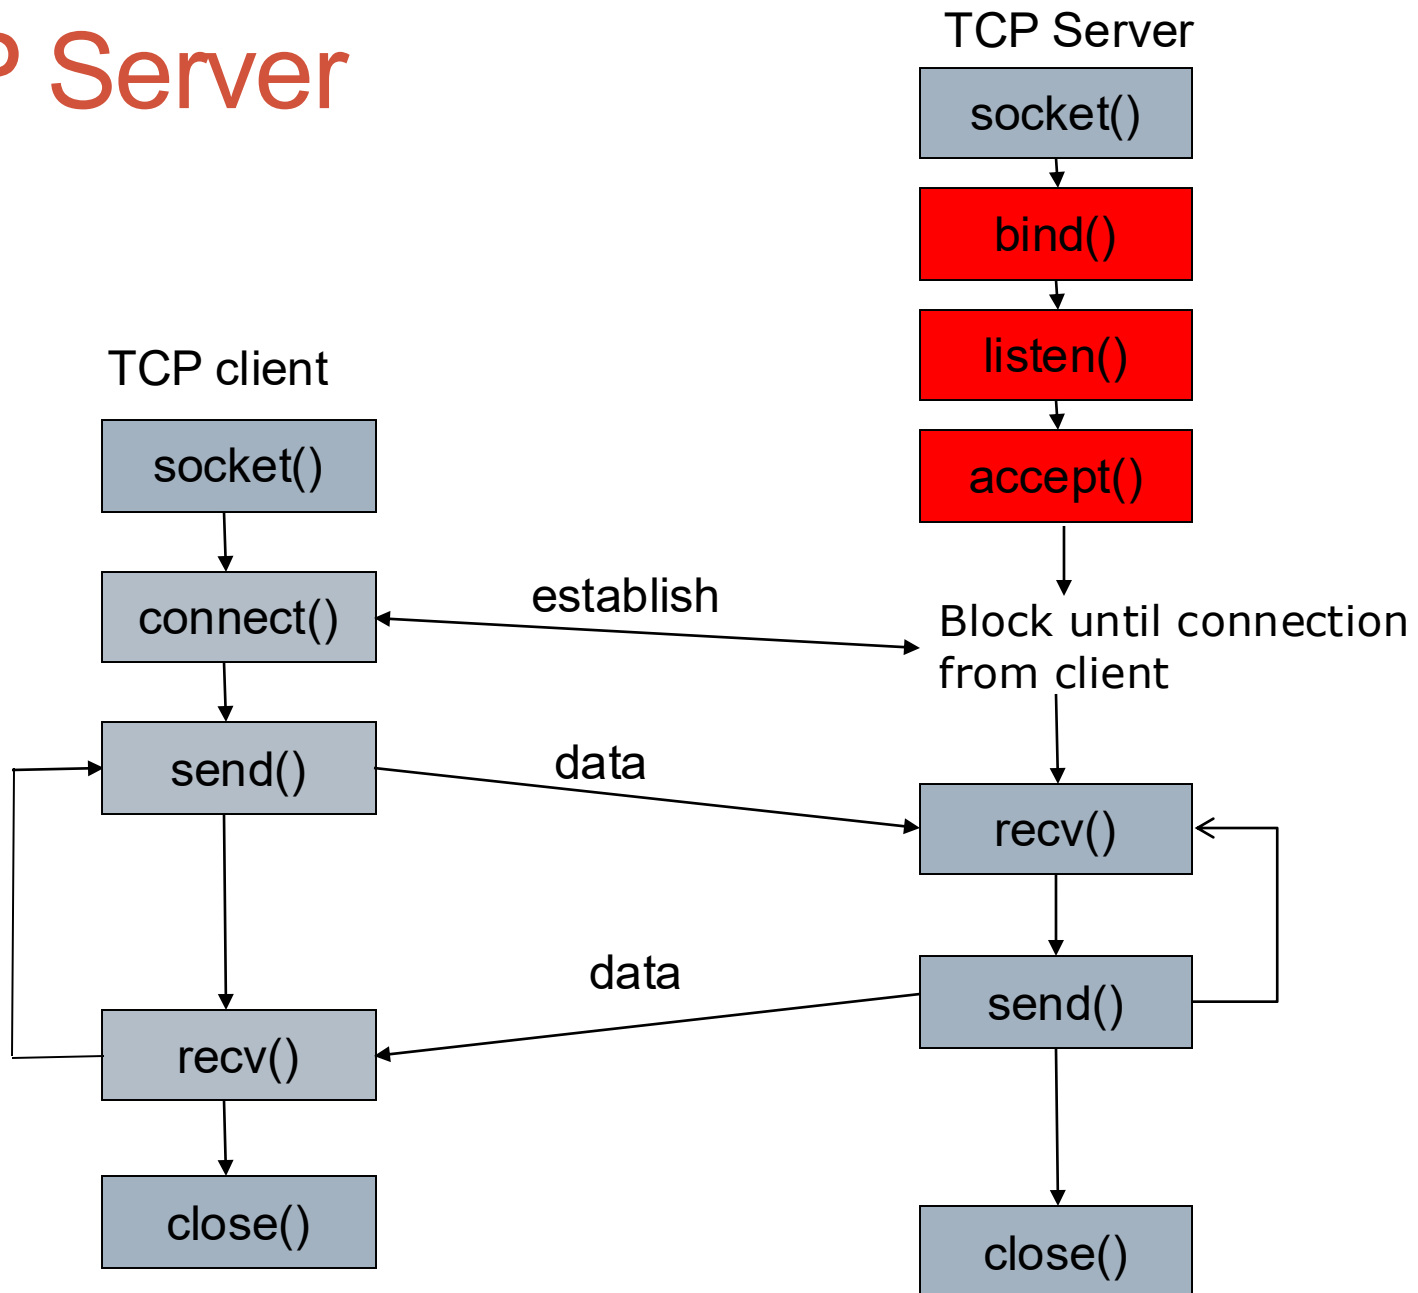
TCP (Transmission Control Protocol)

- Provide reliable communication
- Data rate control
- Example
 - Mail
 - WEB
 - Image




TCP SERVER

TCP Server



TCP server side

1. Create a socket – `socket()`.
 2. Bind the socket – `bind()`.
 3. Listen on the socket – `listen()`.
 4. Accept a connection – `accept()`.
 5. Send and receive data – `recv()`, `send()`.
 6. Disconnect connection– `close()`
 7. Close LISTENING socket
- 
- repeatedly

Socket Mode

- Types of server sockets
 - *Iterating server*: Only one socket is opened at a time.
 - *Concurent server*: After an accept, a child process/thread is spawned to handle the connection.
 - *Multiplexing server*: use select to simultaneously wait on all open socketIds, and waking up the process only when new data arrives

socket ()

```
#include <sys/types.h>
#include <sys/socket.h>
int socket(int domain, int type, int protocol);
```

- Creates an endpoint for communication
- [IN] `domain`: `AF_INET`, `AF_INET6`, ...
- [IN] `type` argument can be:
 - `SOCK_STREAM`: Provides sequenced, reliable, two-way, connection-based byte streams
 - `SOCK_DGRAM`: Supports datagrams
 - `SOCK_RAW`: Provides raw network protocol access
- [IN] `protocol` is usually 0
- Returns value
 - A new socket descriptor that you can use to do sockety things with
 - If error occurs, return -1

bind()

```
#include <sys/types.h>
#include <sys/socket.h>
int bind(int sockfd, const struct sockaddr *addr,
         socklen_t addrlen);
```

- Associate a socket with an IP address and port number
- Where
 - [IN] `sockfd` : socket descriptor
 - [IN] `addr` : pointer to a `sockaddr` structure assigned to `sockfd`
 - [IN] `addrlen` : specifies the size, in bytes of address structure pointed to by `addr`
- Return value
 - Returns 0 if no error occurs.
 - Otherwise, return -1 (and **errno** will be set accordingly)

listen()

```
#include <sys/socket.h>
int listen(int sockfd, int backlog);
```

- Establish a socket to LISTENING for incoming connection.
- Parameters:
 - [IN] `sockfd`: a descriptor identifying a bound, unconnected socket
 - [IN] `backlog`: the number of pending connections the queue will hold
- Return value
 - On success, 0 is returned
 - On error, -1 is returned

accept()

```
#include <sys/types.h>
#include <sys/socket.h>
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

- Accept an incoming connection on a LISTENING socket
- Parameters:
 - `sockfd`: A descriptor identifying a socket which is listening for connections after a `listen()`.
 - `addr`: pointer to a `sockaddr` structure filled in with the address of the peer socket
 - `addrlen`: return the actual size of the peer address.
- Return value
 - Newly connected socket descriptor if no errors
 - -1 if has errors

send()

```
#include <sys/types.h>
#include <sys/socket.h>
ssize_t send(int sockfd, const void *buf, size_t len,
             int flags);
```

- Send data on a connected socket
- Parameter:
 - [IN] `sockfd`: a descriptor identifying a connected socket.
 - [IN] `buf`: points to the buffer containing the message to send.
 - [IN] `len`: specifies the length of the message
 - [IN] `flags`: specifies the type of message transmission, usually 0
- Return value:
 - If no error occurs, `send()` returns the total number of characters sent
 - Otherwise, return -1

send () - Flags

- `MSG_OOB`: Send as “out of band” data. The receiver will receive the signal `SIGURG` and it can then receive this data without first receiving all the rest of the normal data in the queue.
- `MSG_DONTROUTE`: Don't send this data over a router, just keep it local.
- `MSG_DONTWAIT`: If **`send()`** would block because outbound traffic is clogged, have it return `EAGAIN`.
- `MSG_NOSIGNAL`: If you **`send()`** to a remote host which is no longer **`recv()`**, you'll typically get the signal `SIGPIPE`. Adding this flag prevents that signal from being raised.

send()

```
char sendBuff[2048];
int  dataLength, nLeft, idx;

// Fill sendbuff with 2048 bytes of data
nLeft = dataLength;
idx = 0;

while (nLeft > 0){
    // Assume s is a valid, connected stream socket
    ret = send(s, &sendBuff[idx], nLeft, 0);
    if (ret == -1)
    {
        // Error handler
    }
    nLeft -= ret;
    idx += ret;
}
```

recv()

```
#include <sys/types.h>
#include <sys/socket.h>
ssize_t recv(int sockfd, void *buf, size_t len, int flags);
```

- Receive data on a socket
- Parameter:
 - [IN] `sockfd`: a descriptor identifying a connected socket.
 - [IN, OUT] `buf`: points to a buffer where the message should be stored
 - [IN] `len`: specifies the length in bytes of the buffer
 - [IN] `flags`: specifies the type of message reception, usually 0
- Return value:
 - If no error occurs, returns the length of received message in bytes
 - If peer has performed an orderly shutdown, return 0
 - Otherwise, return -1

receive()

```
char    recvBuff[1024];
int     ret, nLeft, idx;
nLeft = dataLength; //length of the data needs to be
                    //received

idx = 0;

while (nLeft > 0)
{
    ret = recv(s, &recvBuff[idx], nLeft, 0);
    if (ret == -1)
    {
        // Error handler
    }
    idx += ret;
    nLeft -= ret;
}
```

recv() - Flags

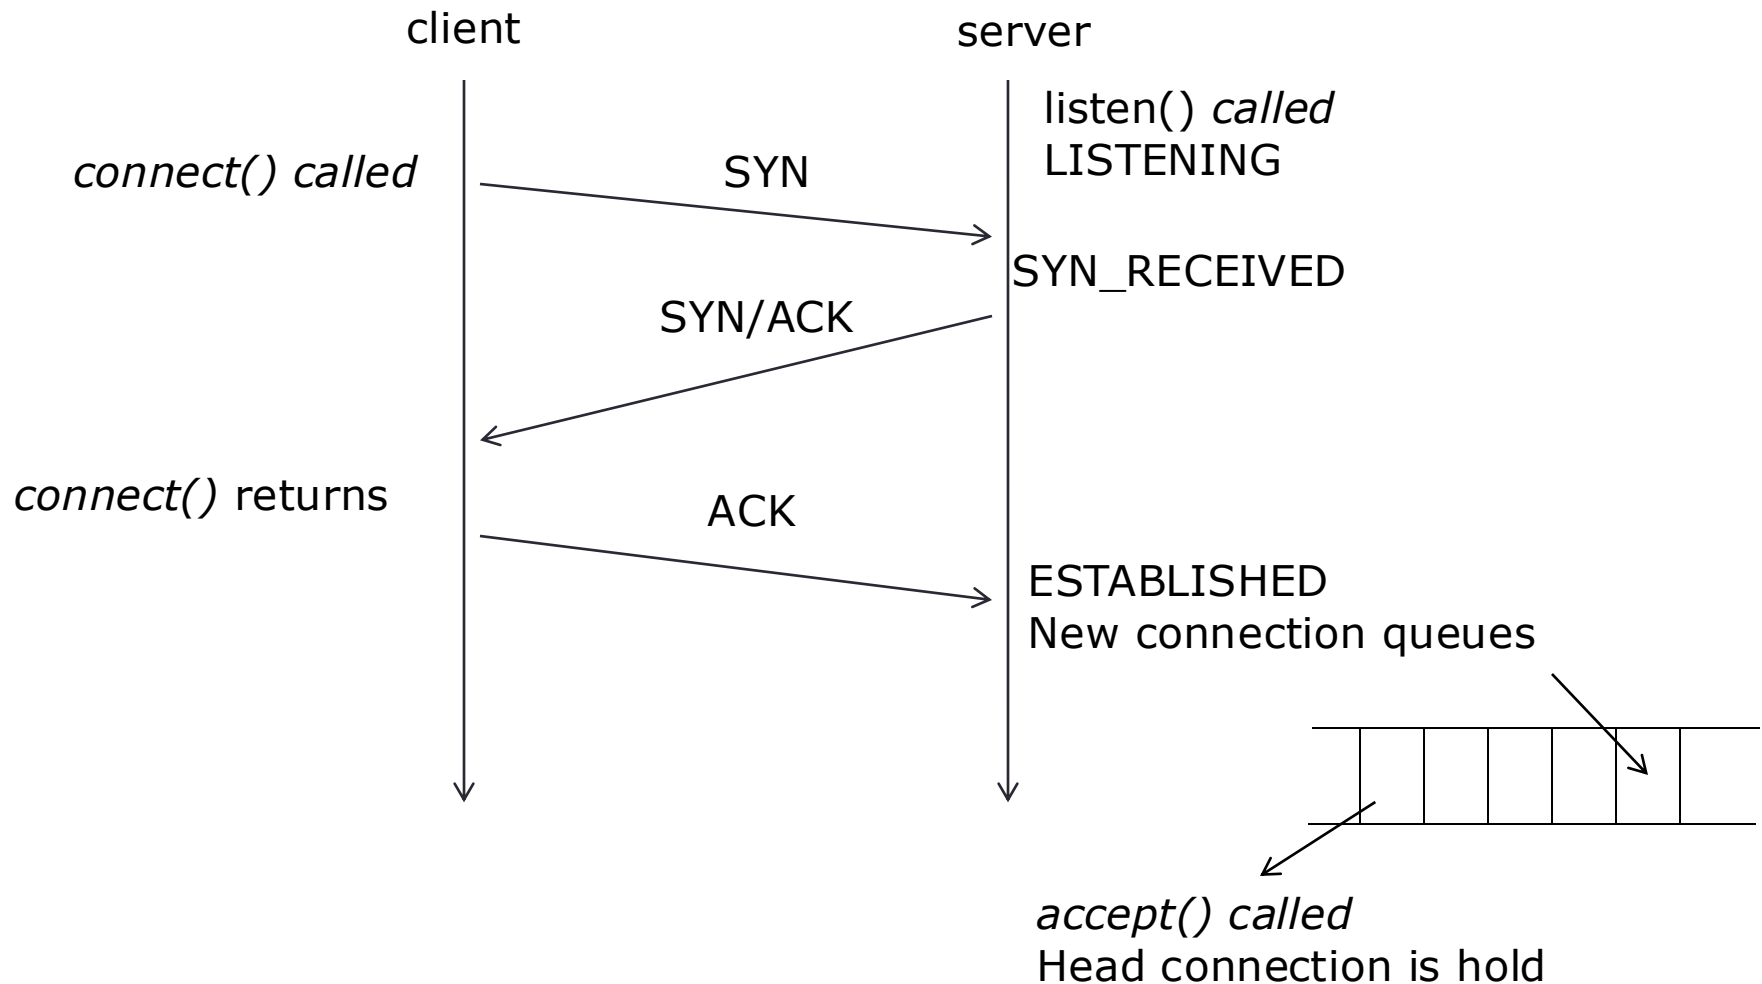
- MSG_PEEK: Peeks at an incoming message. The data is treated as unread and the next *recvfrom()* or similar function shall still return this data.
- MSG_OOB: Requests out-of-band data. The significance and semantics of out-of-band data are protocol-specific.
- MSG_WAITALL: On SOCK_STREAM sockets this requests that the function block until the full amount of data can be returned, excepting:
 - the connection is terminated
 - MSG_PEEK was specified
 - an error is pending for the socket
 - a signal is caught
- Use bitwise OR operator (|) to combine more than one flag

close()

```
#include <unistd.h>
int close(int sockfd);
```

- Close a socket descriptor
- [IN] `sockfd`: a descriptor identifying a socket.
- Return value
 - Returns 0 if no error occurs.
 - Otherwise, return -1 (and **errno** will be set accordingly)
- `close()` **VS** `shutdown()`
 - `close()` tries to complete this transmission before closing, frees the socket descriptor
 - `shutdown()`: immediately stops receiving and transmitting data, don't releases the socket descriptor

Process connections



Example

```
int listenfd, connfd, n;  
pid_t childpid;  
socklen_t clien;  
char buf[MAXLINE];  
struct sockaddr_in cliaddr, servaddr;  
listenfd = socket (AF_INET, SOCK_STREAM, 0);
```

*creation of
server socket*

```
servaddr.sin_family = AF_INET;  
servaddr.sin_addr.s_addr = htonl(INADDR_ANY);  
servaddr.sin_port = htons(SERV_PORT);
```

*Preparation of
the socket
address struct*

```
bind (listenfd, (struct sockaddr *) &servaddr, sizeof(servaddr));
```

*Bind the socket
to the port in
address*

```
listen (listenfd, LISTENQ);
```

*Listen for connection
to the socket*

```
printf("%s\n", "Server running...waiting for connections.");
```

Example (Cont.)

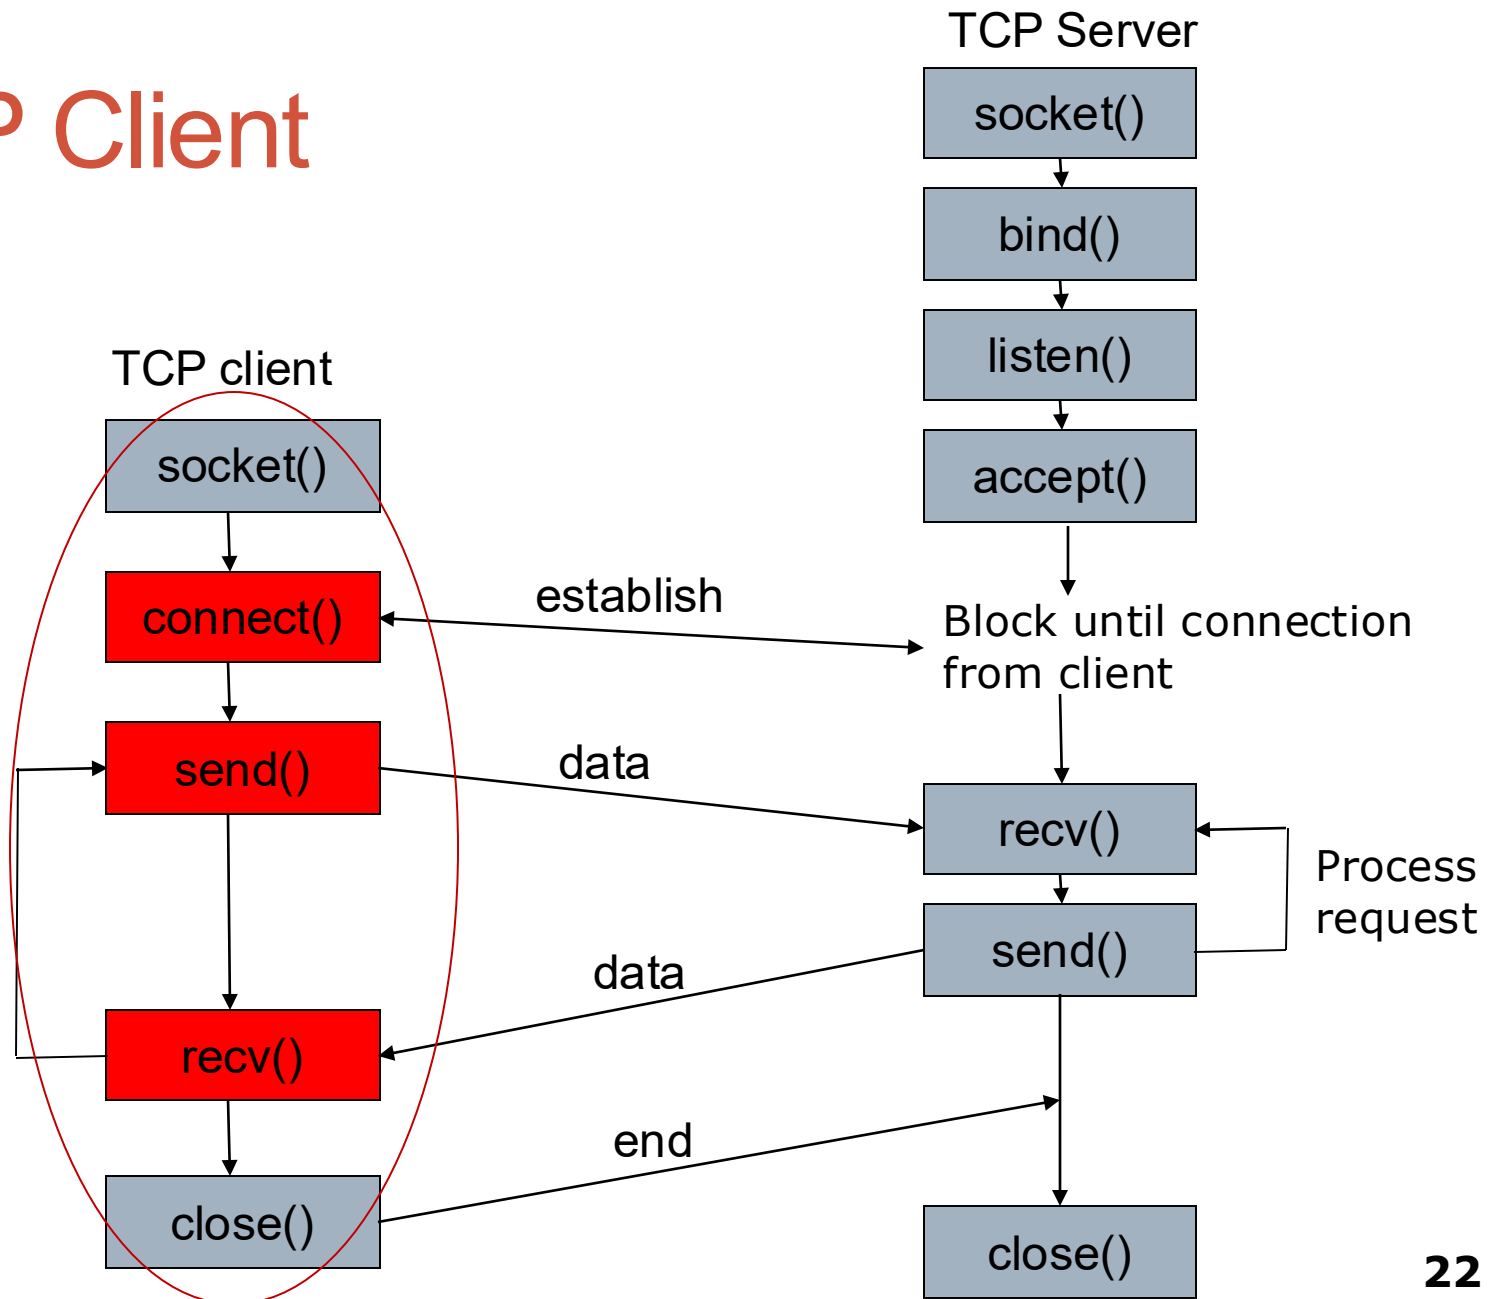
```
for ( ; ; ) {  
    clien = sizeof(cliaddr);  
    connfd = accept (listenfd, (struct sockaddr *) &cliaddr, &clilen);  
    printf("%s\n", "Received request...");  
    while ( (n = recv(connfd, buf, MAXLINE, 0)) > 0) {  
        printf("%s", "String received from and resent to the client:");  
        puts(buf);  
        send(connfd, buf, n, 0);  
    }  
    if (n < 0) {  
        perror("Read error");  
        exit(1);  
    }  
    close(connfd); // close the file descriptor.  
}  
close (listenfd); //close listening socket
```

*Accept a connection
request → return a
File Descriptor (FD)*

*Send and receive
data from the FD*

TCP CLIENT

TCP Client



TCP client side

- The typical TCP client's communication involves four basic steps:
 - Create a TCP socket using `socket()`.
 - Establish a connection to the server using `connect()`.
 - Communicate using `send()` and `recv()`.
 - Close the connection with `close()`.
- Why “clients” doesn't need `bind()` ?

connect ()

```
#include <sys/types.h>
#include <sys/socket.h>
int connect(int sockfd, const struct sockaddr *serv_addr,
            socklen_t addrlen);
```

- Connect a socket to a server
- Parameters:
 - [IN] sockfd: A descriptor identifying an unconnected socket.
 - [IN] serv_addr: The address of the server to which the socket is to be connected.
 - [IN] addrlen: The length of the name.
- Return value
 - If no error occurs, returns 0.
 - Otherwise, it returns -1

send(), receive()

- Similar in TCP server

Example

```
int main(int argc, char **argv)
{ int sockfd;
  struct sockaddr_in servaddr;
  char sendline[MAXLINE], recvline[MAXLINE];
  //Create a socket for the client
  if ((sockfd = socket (AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("Problem in creating the socket");
    exit(2);
  }
  //Creation of the remote server socket information structure
  memset(&servaddr, 0, sizeof(servaddr));
  servaddr.sin_family = AF_INET;
  servaddr.sin_addr.s_addr = inet_addr(argv[1]);
  servaddr.sin_port = htons(SERV_PORT); //convert to big-endian order
```

create a client socket

create a socket
addr info pointing
to server socket

Example (Cont.)

```
// Connect the client to the server socket
```

```
if (connect(sockfd, (struct sockaddr *) &servaddr, sizeof(servaddr))<0) {  
    perror("Problem in connecting to the server");  
    exit(3);  
}
```

```
while (fgets(sendline, MAXLINE, stdin) != NULL) {
```

```
    send(sockfd, sendline, strlen(sendline), 0);
```

```
    if (recv(sockfd, recvline, MAXLINE, 0) == 0) {
```

```
        //error: server terminated prematurely
```

```
        perror("The server terminated prematurely");
```

```
        exit(4);
```

```
    }
```

Connect the client
socket with remote
server

Send and receive
data from client
socket

OTHERS

shutdown ()

```
#include <sys/socket.h>
int shutdown(int socket, int how);
```

- Shut down socket send and receive operations
- Where
 - [IN] sockfd: a descriptor identifying a socket.
 - [IN] how: SHUT_RD, SHUT_WR, SHUT_RDWR
- Return value
 - Returns 0 if no error occurs.
 - Otherwise, return -1

Socket options

```
#include <sys/socket.h>
int setsockopt (int sockfd, int level, int optname,
               void *optval, int optlen);
```

- Set the options that control the transferring data on a socket
- Parameters:
 - [IN] `sockfd`: refer to an open socket descriptor
 - [IN] `level`: specifies the protocol level at which the option resides
 - [IN] `optname`: specifies a single option to set
 - [IN] `optval`: points to the setted option value
 - [IN] `optlen`: the size of option value pointed by `optval`
- Return:
 - Returns 0 if no error occurs.
 - Otherwise, return -1 (and **errno** will be set accordingly)

Socket options(cont)

```
#include <sys/socket.h>
int getsockopt (int sockfd, int level, int optname,
               void *optval, int *optlen);
```

- Set the options that control the transferring data on a socket
- Parameters:
 - [IN] `sockfd`: refer to an open socket descriptor
 - [IN] `level`: specifies the protocol level at which the option resides
 - [IN] `optname`: specifies a single option to set
 - [OUT] `optval`: points to the setted option value
 - [IN, OUT] `optlen`: the size of option value pointed by `optval`
- Return:
 - Returns 0 if no error occurs.
 - Otherwise, return -1 (and **errno** will be set accordingly)

level = SOL_SOCKET

Value name	Type	Description
SO_BROADCAST	int	Configures a socket for sending broadcast data.(Only UDP socket)
SO_DONTROUTE	int	Sets whether outgoing data should be sent on interface the socket is bound to and not a routed on some other interface
SO_KEEPALIVE	int	TCP automatically sends a keep-alive probe to the peer
SO_LINGER	linger	specifies how the close function operates for a connection-oriented protocol
SO_REUSEADDR	int	Allows the socket to be bound to an address that is already in use
SO_RCVTIMEO	timeval	Sets the timeout for blocking receive calls
SO_SNDTIMEO	timeval	Sets the timeout for blocking send calls